



Elder Mind

Agentic Governance

by Cyber Elders

Local-first agentic governance for coding agents

A deterministic pre-tool-use gate that blocks risky actions, checks dependency installs against known-compromised packages, surfaces threat patterns, and can escalate high-risk calls to a multi-model council — on **your** machine, with **your** model.

THE WEDGE

Governance that runs as your coding agent's own pre-tool-use hook — local, deterministic, and auditable — not a control plane your agent phones home to. It governs the exact moment an agent decides to run `rm -rf`, `force-push`, or `curl | bash`.

WHAT'S IN THE HARNESS

- **Pre-tool gate** — $\text{impact} \times \text{likelihood} \rightarrow \text{allow} / \text{warn} / \text{ask} / \text{block}$ on a versioned policy.
- **Supply-chain (opt-in)** — OSV + OpenSSF malicious-packages, offline fallback, release-age.
- **Threat detectors** — MITRE-tagged regex surfacing over tool arguments.
- **Council review (BYO-LLM)** — multi-model deliberation using your own model.
- **Tool-descriptor pinning** — catches MCP/tool 'rug-pulls' after approval.
- **Tamper-evident audit** — hash-chained JSONL; `eldermind verify`.

HONEST STANDARDS POSTURE

OWASP Agentic Top 10 (2026)—**aware** and **aligned** to the NIST AI RMF structure — never 'compliant', 'certified', or 'covers all 10'. Enforces ASI02 / ASI04 / ASI05, audits ASI03, and openly cedes what a local hook structurally can't see.

Govern the action. Teach the operator. Keep it local.