



# Elder Mind

## Agentic Governance

by Cyber Elders

## Local-first agentic governance for coding agents

Your AI coding agent acts on its own — running commands, editing files, installing packages. Elder Mind adds the pause before the one action that could hurt: it blocks it, explains why in plain language, and logs it.

A deterministic pre-tool-use gate that blocks risky actions, checks dependency installs against known-compromised packages, surfaces threat patterns, and can escalate high-risk calls to a multi-model council — on **your** machine, with **your** model.

### THE WEDGE

Governance that runs as your coding agent's own pre-tool-use hook — local, deterministic, and auditable — not a control plane your agent phones home to. It governs the exact moment an agent decides to run `rm -rf`, `force-push`, or `curl | bash`.

### WHAT'S IN THE HARNESS

- **Pre-tool gate** —  $\text{impact} \times \text{likelihood} \rightarrow \text{allow} / \text{warn} / \text{ask} / \text{block}$  on a versioned policy.
- **Supply-chain (opt-in)** — OSV + OpenSSF malicious-packages, offline fallback, release-age.
- **Threat detectors** — MITRE-tagged regex surfacing over tool arguments.
- **Council review (BYO-LLM)** — multi-model deliberation using your own model.
- **Tool-descriptor pinning** — catches MCP/tool 'rug-pulls' after approval.
- **Tamper-evident audit** — hash-chained JSONL; `eldermind verify`.

### HONEST STANDARDS POSTURE

OWASP Agentic Top 10 (2026)—**aware** and **aligned** to the NIST AI RMF structure — never 'compliant', 'certified', or 'covers all 10'. Enforces ASI02 / ASI04 / ASI05, audits ASI03, and openly cedes what a local

hook structurally can't see.

**Govern the action. Teach the operator. Keep it local.**